

CS387 Lab 1 - Simulate a DB server in Python

Your task is to write a Python application that will process data in a set of files given to you (you may hardcode these file names), accept, parse and answer some queries that users supply. The set of files corresponds to data that can potentially be stored in a DB but is stored as CSV files instead and your application effectively behaves like a DB server. File names correspond to relations and headers within the files correspond to column names. Your application should read all of these files on startup and support the following queries where the values are parametrized by user input.

Assume that all keywords (those in bold font) are supplied in lowercase only.

Query types and formats:

NOTE: The format is designed to be simple (expected tokens are provided in file to be submitted), so that you *do not* spend time writing sophisticated parsing code. If you are spending too much time writing parsing code, you are doing something wrong - talk to instructor/TA. You should be able to manage the parsing with just using python string split once or twice in each query.

1. **Query type 1:** Retrieve data from a “relation” - each relation’s data is captured in a single file. This can take one of 3 forms.

a. **select * from** <relation> ;

// outputs values of all columns for all rows

// parsing note: see the space before the semicolon - this is to enable easy parsing

b. **select * from** <relation> **where** <column> = <value> ;

// support only equality

// outputs values of all columns in the relation for a subset of rows defined by the where condition

// parsing note: see the space before the semicolon - this is to enable easy parsing

c. **select** <col1>,<col2>...,<coln> **from** <relation> **where** <colX> = <value> ;

// outputs selected columns in the relation for a subset of rows defined by the where condition. colX is a column of the relation but may not correspond to the columns being output

// parsing note: see the space before the semicolon - this is to enable easy parsing

// parsing note: also note that there is *no* space between different column names - you can first split using space, then split using comma

2. **Query type 2:** Join data from 2 tables

select * from <relation1>,<relation2> **where** <relation1>.<column1> = <relation2>.<column2> ;

// outputs all columns of both relations for those rows in the cross product of the two relations where the condition specified is satisfied

// parsing note: see the space before the semicolon - this is to enable easy parsing

// parsing note: also note that there is *no* space between the two relation names - you can first split using space, then split using comma

3. Query type 3: Aggregates

```
select count from <relation> where <column> = <value> ;
```

// outputs an integer value of the count of rows in the relation where the condition is satisfied.

// parsing note: see the space before the semicolon - this is to enable easy parsing

4. Query type 4: Set operations (intersection and union only)

```
a. ( select <column> from <relation1> where <colX> = <value1> ) intersect  
( select <column> from <relation2> where <colY> = <value2> ) ;
```

b. As above, but **intersect** replace intersect by **union**

// outputs the set of all values of <column> from rows in <relation1> where the condition is satisfied on <relation1> INTERSECTED/UNIONED with the set of values of <column> from the rows in <relation2> where the condition is satisfied on relation2. Note that <column> should be the same across both relations.

// parsing note: see the space before the semicolon - this is to enable easy parsing

// parsing note: see space before after the open/close brackets - this is to enable easy parsing

Formatting your output: For each row, output the columns in the order supplied in the CSV files comma separated with NO additional spaces anywhere. Each row must be on a new line. In the case of joins output all the columns of the first relation followed by the columns of the second relation in that order. In the case of union output in the order supplied in the file of first relation followed by that of second relation.

See sample outputs given for the example database.

The Interface: Executing your application

Upon running the python app, it should take the input in the format below in a while loop and return the results . The Query Type = 0 is a signal to exit the application. Text in bold is input.

```
~$ python3 a1.py
```

```
Query Type? 1b
```

```
Enter your query: select * from student where dept_name = Physics ;
```

```
<output>
```

```
Query Type? 3
```

```
Enter your query: select count from course where credits = 4 ;
```

```
<output>
```

```
.
```

```
.
```

```
output
```

```
Query Type? 0
```

```
exiting...
```

```
~$
```

Set of data files included

Each file is a csv with a header that corresponds to column names. The name of the file is the name of the relation. Follow `labDirectory` structure in VLab and **do not** use absolute path for csv files.

`advisor, classroom, course, department, instructor, prereq, section, student, takes, teaches, time_slot`

Testing your application

We have included sample test case for each type of query as `input.txt` and expected output as `output.txt`. Note that we will be testing with cases other than these (but conforming to the same structure) for grading.

Submission Instructions

Submit a single Python file named `a1.py` (file with expected tokens is provided in `submission` directory) via VLab.

Grading Rubric

Query Type	Marks
1a	5
1b	15
1c	15
2	25
3	10
4a	15
4b	15
Total	100